


Quantum Information and Computing

Academic Year 2024 - 2025

ASSIGNMENT 1

Due for October 28, 2024



Gabriel Amorosetti

gabriel.amorosetti@studenti.unipd.it

1. Setup

- Creation of a working directory on Windows through the terminal

```
1 PS C:\Users\Gabriel\Documents\Physics of Data\2nd Year\Quantum Information and Computing> mkdir assignment1
```

- Submission of a test job in Fortran :
 - Visual Studio Code 1.94.2
 - Modern Fortran extension 3.2.0
 - GNU Fortran 14.1.0 compiler

- Compile with gfortran :

```
1 gfortran -o a1_ex1.exe a1_ex1.f90
```

```
1 program Setup
2   implicit none ! To enforce explicit variable declaration
3   integer :: a,b
4   real :: x, r
5
6   ! Simple calculation as a test job
7   x = 3.2
8   a = 3
9   b = 4
10  r = b*x - a
11
12  print *, "Hello world, here is a simple calculation : 4*3.2 - 3 = ", r
13
14 end program Setup
```

- Output :

```
1 PS C:\Users\Gabriel\Documents\Physics of Data\2nd Year\Quantum Information and Computing\assignment1> .\a1_ex1.exe
2 Hello world, here is a simple calculation : 4*3.2 - 3 = 9.80000019
```

2. Number precision

- Sum 2.000.000 and 1 with INTEGER*2 and INTEGER*4
 - Implemented in the code with
- Sum $\pi \cdot 10^{32}$ and $\sqrt{2} \cdot 10^{21}$ with single and double precision
 - Implemented in the code with

```
1 integer(2) ! 16-bit integers
2 integer(4) ! 32-bit integers
```

```
1 real(4) ! Single precision real numbers (32-bit)
2 real(8) ! Double precision real numbers (64-bit)
```

- Compilation error :

```
1 .\a1_ex2.f90:20:10:
2
3   20 |   a_2 = 2000000   ! Arithmetic overflow, 2 000 000 is too big for a 16-bit integer (only goes from -32 768 to 32 767)
4       |           1
5 Error: Arithmetic overflow converting INTEGER(4) to INTEGER(2) at (1). This check can be disabled with the option '-fno-range-check'
```

- Forcing compilation with '-fno-range-check' and executing :

```
1 Sum 2 000 000 and 1 with INTEGER*2 : -31615
2 Sum 2 000 000 and 1 with INTEGER*4 : 2000001
3 Sum of pi*10^32 and sqrt(2)*10^21 with single precision : 3.14159278E+32
4 Sum of pi*10^32 and sqrt(2)*10^21 with double precision : 3.1415926536039354E+032
```

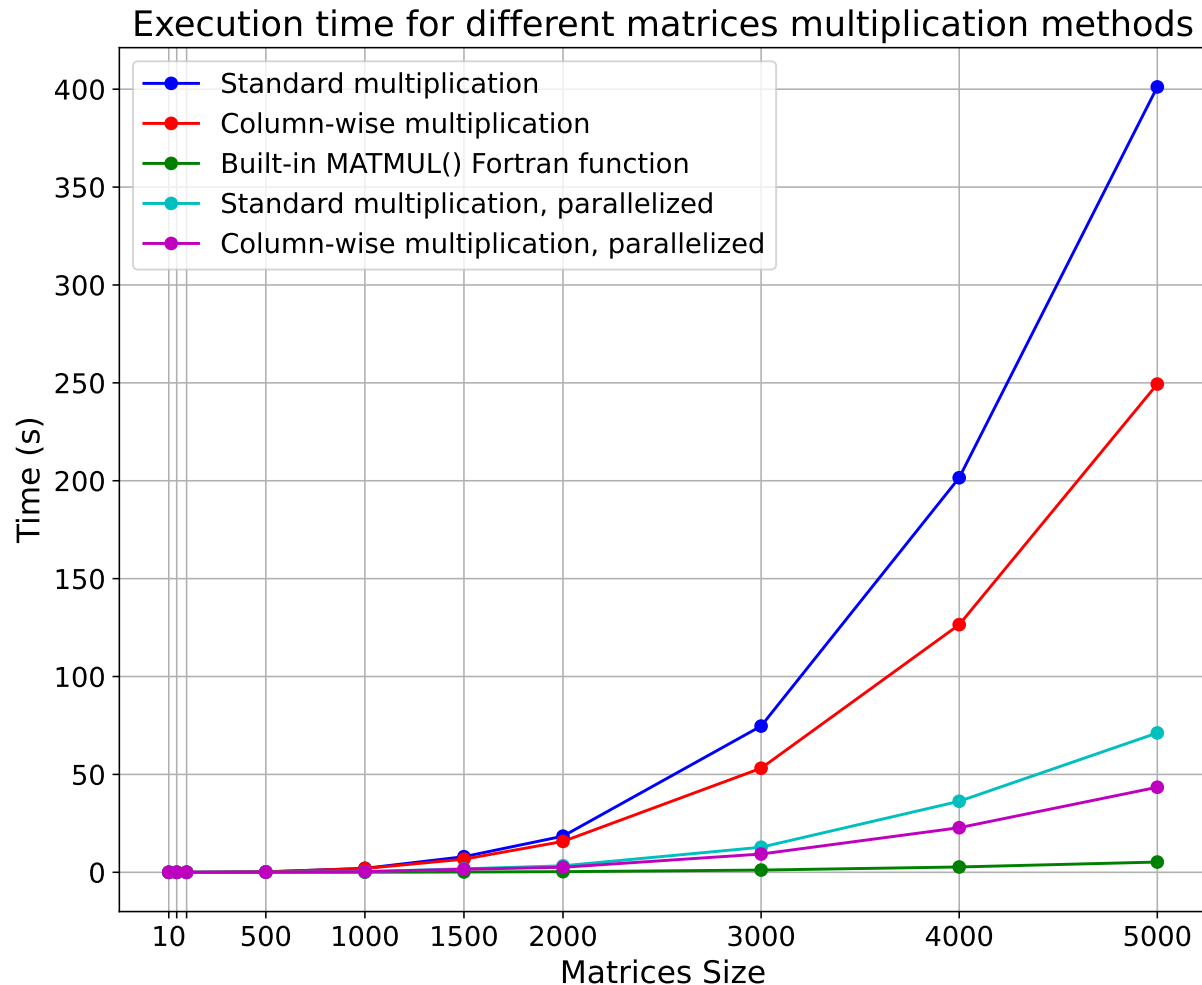
- Limits of 16-bit integers and 32-bit reals :
 - INTEGER*2 can only represent numbers from from -32 768 to 32 767, causing inaccurate results, whereas INTEGER*4 can represent numbers from -2^{31} to $2^{31} - 1$, giving the correct sum
 - Single precision ignores the smaller terms with important different magnitudes in the sum, while Double precision takes in account these smaller details, giving a more precise result

3. Matrix multiplication

- 3 approaches to compute $A \times B = C$
 - Standard multiplication :
iterates through the rows of the A matrix and the columns of the B matrix,
filling the output matrix C by row
 - Column-wise multiplication :
iterates through the columns of the A and B matrices,
filling the output matrix C by column
 - Assumed to be faster as Fortran stores arrays in column-major order
 - Fortran intrinsic function :

```
1 C = matmul(A,B)
```


3. Matrix multiplication : methods comparison




- Additionally, parallelized versions of the standard and column-wise methods were implemented, in order to use all the CPU cores during computation
- Standard multiplication is the slowest one, while the column-wise multiplication performs better, due to a more efficient memory access in Fortran
- Parallelized versions of these methods improve the execution times, the standard parallelized method being still worse than the column-wise one
- The built-in MATMUL() function is significantly faster than the previous approaches, maintaining low execution time for large sizes, as it is the most optimized and efficient

3. Matrix multiplication : optimization flags


Standard multiplication execution time with different compiler optimization flags




Column-wise multiplication execution time with different compiler optimization flags




MATMUL() Fortran function execution time with different compiler optimization flags



Standard parallelized multiplication execution time with different compiler optimization flags



Column-wise parallelized multiplication execution time with different compiler optimization flags



- The MATMUL() function does not significantly benefit from any optimization flags, as it already includes low-level optimizations
- The -O flags, listed from the least to most aggressive, improve the execution time according to their respective level of code optimization
- -funroll-loops, performing loop unrolling, actually slows the execution
- The combination of different flags seems to be the most efficient approach here
- Further explanations of the different flags are in the code